

ASSIGNMENT MULTIPLE REGRESSION

In [81]:

```

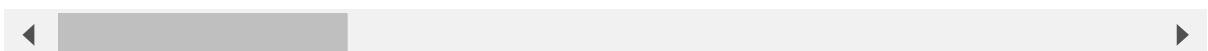
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 from mpl_toolkits.mplot3d import Axes3D
5 from sklearn import linear_model
6 df = pd.read_csv("cars.csv")
7 df

```

Out[81]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Engine Information.Driveline	Information.
0	140	143	202	All-wheel drive	Audi 3.2L 6 250hp 2
1	140	143	202	Front-wheel drive	Audi 2.0L 4 200 hp 20
2	140	143	202	Front-wheel drive	Audi 2.0L 4 200 hp 20
3	140	143	202	All-wheel drive	Audi 2.0L 4 200 hp 20
4	140	143	202	All-wheel drive	Audi 2.0L 4 200 hp 20
...
5071	13	253	201	Front-wheel drive	Honda Cylinder 250
5072	141	249	108	All-wheel drive	Lamborghini cylinder 552
5073	160	249	108	All-wheel drive	Lamborghini cylinder 552
5074	200	210	110	Rear-wheel drive	BMW cylinder 315 ft-lb
5075	200	94	110	Rear-wheel drive	BMW cylinder 315 ft-lb

1076 rows × 18 columns



In [82]: 1 df.describe()

Out[82]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Information.Number of Forward Gears	Engine Informatic
count	5076.000000	5076.000000	5076.000000	5076.000000	5076.000000
mean	145.632191	127.825847	144.012411	5.519110	17.2
std	62.125026	77.358295	79.925899	0.845637	4.4
min	1.000000	2.000000	1.000000	4.000000	8.0
25%	104.000000	60.000000	62.000000	5.000000	14.0
50%	152.000000	128.000000	158.000000	6.000000	17.0
75%	193.000000	198.000000	219.000000	6.000000	20.0
max	255.000000	255.000000	254.000000	8.000000	38.0



In [83]: 1 dimension=df.shape
2 print("DIMENSION=",dimension)

DIMENSION= (5076, 18)

In [84]:

```
1 numerical_columns=df.select_dtypes(include=['number'])
2 count=numerical_columns.shape[1]
3 print(count)
4 numerical_columns
```

9

Out[84]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Information.Number of Forward Gears	Engine
0	140	143	202	6	
1	140	143	202	6	
2	140	143	202	6	
3	140	143	202	6	
4	140	143	202	6	
...
5071	13	253	201	5	
5072	141	249	108	6	
5073	160	249	108	6	
5074	200	210	110	6	
5075	200	94	110	6	

5076 rows × 9 columns



```
In [85]: 1 categorical_columns=df.select_dtypes(include=['object', 'category'])
2 count=categorical_columns.shape
3 print(count)
4 categorical_columns
```

(5076, 8)

Out[85]:

	Engine Information.Driveline	Engine Information.Engine Type	Engine Information.Transmission	Fuel Information.Fuel Type	Identifi Type
0	All-wheel drive	Audi 3.2L 6 cylinder 250hp 236ft-lbs	6 Speed Automatic Select Shift	Gasoline	
1	Front-wheel drive	Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo	6 Speed Automatic Select Shift	Gasoline	
2	Front-wheel drive	Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo	6 Speed Manual	Gasoline	
3	All-wheel drive	Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo	6 Speed Automatic Select Shift	Gasoline	
4	All-wheel drive	Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo	6 Speed Automatic Select Shift	Gasoline	
...
5071	Front-wheel drive	Honda 3.5L 6 Cylinder 250 hp 253 ft-lbs	5 Speed Automatic	Gasoline	
5072	All-wheel drive	Lamborghini 5.2L 10 cylinder 552 hp 398 ft-lbs	6 Speed Manual	Gasoline	
5073	All-wheel drive	Lamborghini 5.2L 10 cylinder 552 hp 398 ft-lbs	6 Speed Manual	Gasoline	
5074	Rear-wheel drive	BMW 3.0L 6 cylinder 315hp 330 ft-lbs Turbo	6 Speed Automatic Select Shift	Gasoline	
5075	Rear-wheel drive	BMW 3.0L 6 cylinder 315hp 330 ft-lbs Turbo	6 Speed Automatic Select Shift	Gasoline	

5076 rows × 8 columns



```
In [86]: 1 missing_values = df.isnull().sum()
2 print("Missing values in each column:\n", missing_values)
3
```

Missing values in each column:

Dimensions.Height	0
Dimensions.Length	0
Dimensions.Width	0
Engine Information.Driveline	0
Engine Information.Engine Type	0
Engine Information.Hybrid	0
Engine Information.Number of Forward Gears	0
Engine Information.Transmission	0
Fuel Information.City mpg	0
Fuel Information.Fuel Type	0
Fuel Information.Highway mpg	0
Identification.Classification	0
Identification.ID	0
Identification.Make	0
Identification.Model Year	0
Identification.Year	0
Engine Information.Engine Statistics.Horsepower	0
Engine Information.Engine Statistics.Torque	0

dtype: int64

```
In [87]: 1 duplicates = df.duplicated()
2 num_duplicates = duplicates.sum()
3 print(f"Number of duplicate rows: {num_duplicates}")
4
5 duplicate_rows = df[duplicates]
6 print("Duplicate rows:\n", duplicate_rows)
7
```

4	2009 Audi A3	200
9		▲
121	2011 Nissan 370Z Coupe	201
1		▼
1389	2010 Chevrolet Corvette Grand Sport Convertible	201
0		▼
2203	2011 GMC Canyon	201
1		▼
2343	2011 GMC Terrain	201
1		▼
3569	2011 Toyota Tundra	201
1		▼
3570	2011 Toyota Tundra	201
1		▼
3640	2010 Chevrolet Colorado	201
0		▼
3703	2011 Ford Ranger	201
1		▼
3704	2010 Mercedes-Benz R-Class	201
1		▼

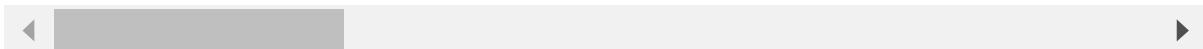
```
In [88]: 1 # Remove duplicate rows and keep the first occurrence
          2 df = df.drop_duplicates()
```

```
In [89]: 1 df
```

Out[89]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Engine Information.Driveline	Information
0	140	143	202	All-wheel drive	Audi 3.2L 250hp
1	140	143	202	Front-wheel drive	Audi 2.0L 200 hp
2	140	143	202	Front-wheel drive	Audi 2.0L 200 hp
3	140	143	202	All-wheel drive	Audi 2.0L 200 hp
5	91	17	62	All-wheel drive	Audi 3.2L 265hp
...
5071	13	253	201	Front-wheel drive	Hor Cylinder 2
5072	141	249	108	All-wheel drive	Lamborghini cylinder 5
5073	160	249	108	All-wheel drive	Lamborghini cylinder 5
5074	200	210	110	Rear-wheel drive	BMW cylinder 3 ft.
5075	200	94	110	Rear-wheel drive	BMW cylinder 3 ft.

5058 rows × 18 columns



In [90]: 1 df.dtypes

```
Out[90]: Dimensions.Height           int64
Dimensions.Length            int64
Dimensions.Width             int64
Engine Information.Driveline object
Engine Information.Engine Type object
Engine Information.Hybrid      bool
Engine Information.Number of Forward Gears   int64
Engine Information.Transmission    object
Fuel Information.City mpg        int64
Fuel Information.Fuel Type       object
Fuel Information.Highway mpg     int64
Identification.Classification   object
Identification.ID               object
Identification.Make              object
Identification.Model Year        object
Identification.Year              int64
Engine Information.Engine Statistics.Horsepower  int64
Engine Information.Engine Statistics.Torque        int64
dtype: object
```

In [91]: 1 num_cols=['Dimensions.Height','Dimensions.Length','Dimensions.Width',
2 'Engine Information.Number of Forward Gears','Fuel Information.
3 'Fuel Information.Highway mpg','Identification.Year','Engine Inf
4 'Engine Information.Engine Statistics.Torque']
5
6 cat_cols=['Engine Information.Driveline','Engine Information.Engine Type'
7 'Engine Information.Transmission','Fuel Information.Fuel Type',
8 'Identification.Classification','Identification.ID','Identificat
9
10 boolean_cols=['Engine Information.Hybrid']

In [92]: 1 df

Out[92]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Engine Information.Driveline	Inform
0	140	143	202	All-wheel drive	Audi 2
1	140	143	202	Front-wheel drive	Audi 20
2	140	143	202	Front-wheel drive	Audi 20
3	140	143	202	All-wheel drive	Audi 20
5	91	17	62	All-wheel drive	Audi 20

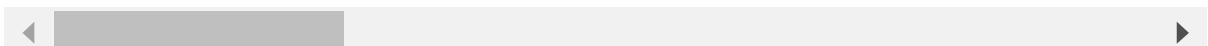
```
In [93]: 1 # import seaborn as sns
2 # sns.histplot(df[num_cols].values,kde=True)
```

```
In [94]: 1 #Checking outliers
2 for i in num_cols:
3     q1=df[i].quantile(0.25)
4     q3=df[i].quantile(0.75)
5     iqr=q3-q1
6     lowerbound=q1-1.5*iqr
7     upperbound=q3+1.5*iqr
8     df=df[(df[i]>=lowerbound) & (df[i]<=upperbound)]
9 df# outlier rows removed
```

Out[94]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Engine Information.Driveline	Informatio
0	140	143	202	All-wheel drive	Audi 3.2L 250hp
1	140	143	202	Front-wheel drive	Audi 2.0L 200 hp
2	140	143	202	Front-wheel drive	Audi 2.0L 200 hp
3	140	143	202	All-wheel drive	Audi 2.0L 200 hp
5	91	17	62	All-wheel drive	Audi 3.2L 265hp
...
5069	3	253	201	Front-wheel drive	Hor Cylinder 2:
5070	3	253	201	Four-wheel drive	Hor Cylinder 2:
5071	13	253	201	Front-wheel drive	Hor Cylinder 2:
5074	200	210	110	Rear-wheel drive	BN cylinder 3 ft.
5075	200	94	110	Rear-wheel drive	BN cylinder 3 ft.

4794 rows × 18 columns



In [95]: 1 df.nunique()

```
Out[95]: Dimensions.Height          189
Dimensions.Length           202
Dimensions.Width            130
Engine Information.Driveline      4
Engine Information.Engine Type    468
Engine Information.Hybrid        1
Engine Information.Number of Forward Gears 4
Engine Information.Transmission   10
Fuel Information.City mpg        22
Fuel Information.Fuel Type       4
Fuel Information.Highway mpg     30
Identification.Classification    2
Identification.ID                4768
Identification.Make              43
Identification.Model Year        828
Identification.Year              4
Engine Information.Engine Statistics.Horsepower 188
Engine Information.Engine Statistics.Torque      193
dtype: int64
```

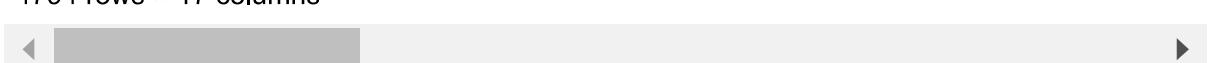
In [96]:

```
1 df=df.drop(columns=['Engine Information.Hybrid'])
2 df
```

Out[96]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Engine Information.Driveline	Information
0	140	143	202	All-wheel drive	Audi 3.2L 250hp
1	140	143	202	Front-wheel drive	Audi 2.0L 200 hp
2	140	143	202	Front-wheel drive	Audi 2.0L 200 hp
3	140	143	202	All-wheel drive	Audi 2.0L 200 hp
5	91	17	62	All-wheel drive	Audi 3.2L 265hp
...
5069	3	253	201	Front-wheel drive	Hor Cylinder 2
5070	3	253	201	Four-wheel drive	Hor Cylinder 2
5071	13	253	201	Front-wheel drive	Hor Cylinder 2
5074	200	210	110	Rear-wheel drive	BM cylinder 3 ft
5075	200	94	110	Rear-wheel drive	BM cylinder 3 ft

4794 rows × 17 columns



In [97]:

```
1 df[cat_cols].nunique()
```

Out[97]:

Engine Information.Driveline	4
Engine Information.Engine Type	468
Engine Information.Transmission	10
Fuel Information.Fuel Type	4
Identification.Classification	2
Identification.ID	4768
Identification.Make	43
Identification.Model Year	828
dtype: int64	

In [98]:

```
1 #Now we will encode categorical data .
2 #target encoding for the 4 highest values
3 target='Fuel Information.City mpg'
4 target_columns = [
5     'Engine Information.Engine Type',
6     'Engine Information.Transmission',
7     'Identification.Model Year',
8     'Identification.ID',
9     'Identification.Make'
10 ]
11
12 for i in target_columns:
13     df[i]=df.groupby(i)[target].transform('mean')
14
15 # onehot encoding for Low cardinality features
16
17 onehot_columns = ['Engine Information.Driveline',
18     'Fuel Information.Fuel Type',
19     'Identification.Classification']
20 df=pd.get_dummies(df,columns=onehot_columns)
```

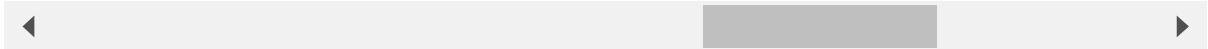
In [104]:

```
1 bool_to_num_cols=['Engine Information.Driveline_All-wheel drive',
2                     'Engine Information.Driveline_Four-wheel drive',
3                     'Engine Information.Driveline_Front-wheel drive',
4                     'Engine Information.Driveline_Rear-wheel drive',
5                     'Fuel Information.Fuel Type_Compressed natural gas',
6                     'Fuel Information.Fuel Type_Diesel fuel',
7                     'Fuel Information.Fuel Type_E85',
8                     'Fuel Information.Fuel Type_Gasoline',
9                     'Identification.Classification_Automatic transmission',
10                    'Identification.Classification_Manual transmission']
11 df[bool_to_num_cols]=df[bool_to_num_cols].astype(int)
```

In [105]: 1 df

Out[105]:

Engine Information.Driveline_Front- wheel drive	Engine Information.Driveline_Rear- wheel drive	Fuel Information.Fuel Type_Compressed natural gas	Fuel Information.Fuel Type_Diesel fuel	Fuel Information.Fuel Type_E85
0	0	0	0	0
1	0	0	0	0
1	0	0	0	0
0	0	0	0	0
0	0	0	0	0
...
1	0	0	0	0
0	0	0	0	0
1	0	0	0	0
0	1	0	0	0
0	1	0	0	0

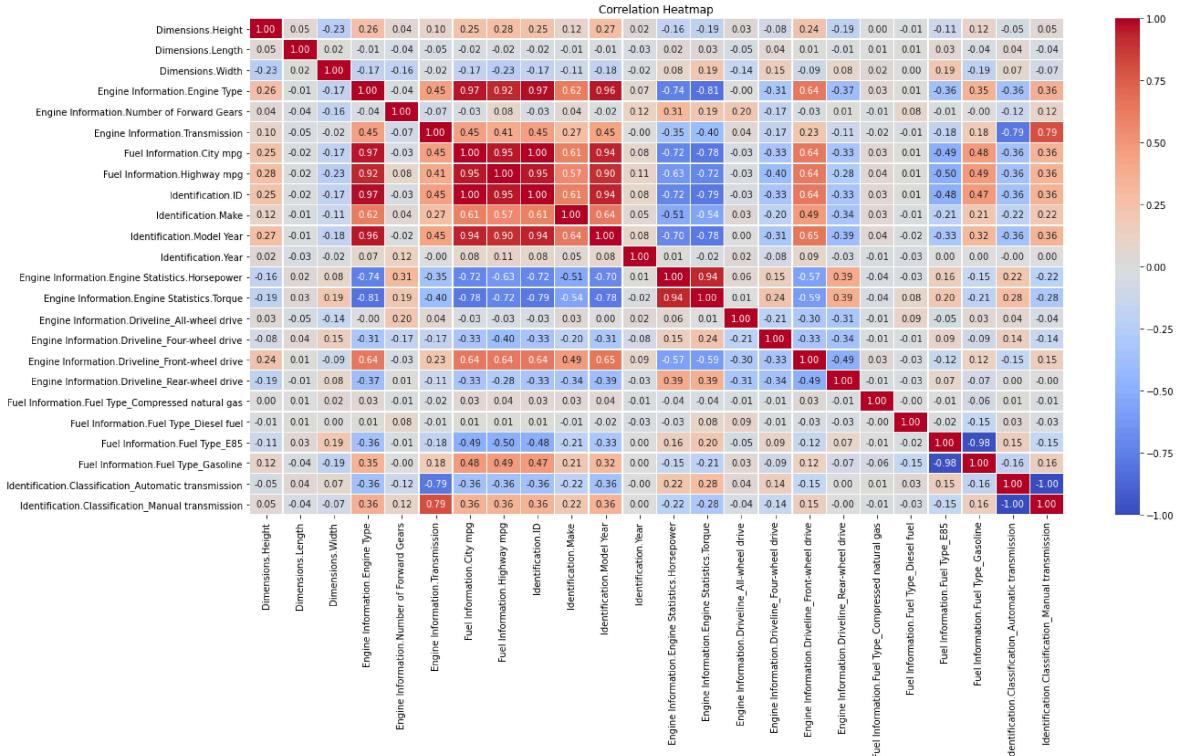


In [109]:

```

1 import seaborn as sns
2 corr_matrix=df.corr()
3 plt.figure(figsize=(20,10))
4 sns.heatmap(corr_matrix,annot=True,cmap='coolwarm',fmt='.2f',linewdiths=0
5 plt.title('Correlation Heatmap')
6 plt.show()
7

```



FEATURE SCALING

In [110]:

```
1 from sklearn.preprocessing import MinMaxScaler
```

In [115]:

```

1 scaler=MinMaxScaler()
2 df[df.select_dtypes(include='number').columns]=scaler.fit_transform(df[df.select_dtypes(include='number').columns])

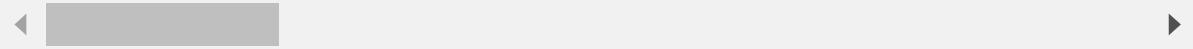
```

In [116]: 1 df

Out[116]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Information.Engine Type	Information.of Forwa
0	0.545455	0.557312	0.794466	0.463415	(
1	0.545455	0.557312	0.794466	0.617886	(
2	0.545455	0.557312	0.794466	0.617886	(
3	0.545455	0.557312	0.794466	0.617886	(
5	0.351779	0.059289	0.241107	0.414634	(
...
5069	0.003953	0.992095	0.790514	0.404878	(
5070	0.003953	0.992095	0.790514	0.404878	(
5071	0.043478	0.992095	0.790514	0.404878	(
5074	0.782609	0.822134	0.430830	0.414634	(
5075	0.782609	0.363636	0.430830	0.414634	(

4794 rows × 24 columns



Feature Selection

In [117]: 1 df.columns

Out[117]: Index(['Dimensions.Height', 'Dimensions.Length', 'Dimensions.Width',
 'Engine Information.Engine Type',
 'Engine Information.Number of Forward Gears',
 'Engine Information.Transmission', 'Fuel Information.City mpg',
 'Fuel Information.Highway mpg', 'Identification.ID',
 'Identification.Make', 'Identification.Model Year',
 'Identification.Year',
 'Engine Information.Engine Statistics.Horsepower',
 'Engine Information.Engine Statistics.Torque',
 'Engine Information.Driveline_All-wheel drive',
 'Engine Information.Driveline_Four-wheel drive',
 'Engine Information.Driveline_Front-wheel drive',
 'Engine Information.Driveline_Rear-wheel drive',
 'Fuel Information.Fuel Type_Compressed natural gas',
 'Fuel Information.Fuel Type_Diesel fuel',
 'Fuel Information.Fuel Type_E85', 'Fuel Information.Fuel Type_Gasolin
 e',
 'Identification.Classification_Automatic transmission',
 'Identification.Classification_Manual transmission'],
 dtype='object')

Using RFE method

In [118]:

```

1 from sklearn.feature_selection import RFE
2 from sklearn.linear_model import Ridge
3
4 X = df.drop('Fuel Information.City mpg', axis=1) # storing independent var
5 y = df['Fuel Information.City mpg'] #target variable
6
7 #creating objects for ridge and rfe
8 ridge = Ridge()
9 rfe = RFE(ridge, n_features_to_select=5) # Select top 5 features
10 rfe.fit(X, y)
11
12 selected_features = X.columns[rfe.support_]
13 print("Selected Features by RFE:", selected_features)

```

Selected Features by RFE: Index(['Engine Information.Engine Type', 'Fuel Information.Highway mpg',
 'Identification.ID', 'Identification.Model Year',
 'Fuel Information.Fuel Type_E85'],
 dtype='object')

In [132]:

```

1 #vif
2 from statsmodels.stats.outliers_influence import variance_inflation_factor
3 selected_features = ['Engine Information.Engine Type', 'Fuel Information.'
4                      'Identification.ID',
5                      'Identification.Model Year',
6                      'Fuel Information.Fuel Type_E85']
7
8 X = df[selected_features]
9
10 # Create a DataFrame to store VIF values
11 vif_data = pd.DataFrame()
12 vif_data["Feature"] = X.columns
13 vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(
14
15 print(vif_data)

```

	Feature	VIF
0	Engine Information.Engine Type	153.036549
1	Fuel Information.Highway mpg	62.119350
2	Identification.ID	179.245052
3	Identification.Model Year	51.423572
4	Fuel Information.Fuel Type_E85	1.282206

In []:

1 #here Fuel Information.Fuel Type_E85 has vif value near to 1, ie, less mu
 2 #So now VIF values of other 4 features should be reduced using PCA

In [133]:

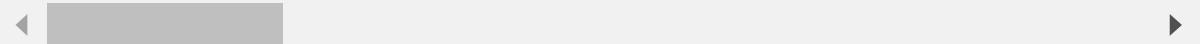
```
1 from sklearn.decomposition import PCA
2 # Extract features with high VIF
3 high_vif_features = ['Engine Information.Engine Type', 'Fuel Information.
4
5
6 # Apply PCA
7 pca = PCA(n_components=1) # Reduce to a single component
8 X_pca = pca.fit_transform(df[high_vif_features])
9
10 # Create a new column in the DataFrame for the PCA component
11 df['PCA_EngineType_highway'] = X_pca
12
13 # Drop the original correlated features
14 df = df.drop(columns=high_vif_features)
15
16 print("Transformed DataFrame with PCA feature added:")
17 df.head()
18
```

Transformed DataFrame with PCA feature added:

Out[133]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Engine Information.Number of Forward Gears	Information.Tr
0	0.545455	0.557312	0.794466	0.666667	
1	0.545455	0.557312	0.794466	0.666667	
2	0.545455	0.557312	0.794466	0.666667	
3	0.545455	0.557312	0.794466	0.666667	
5	0.351779	0.059289	0.241107	0.666667	

5 rows × 23 columns



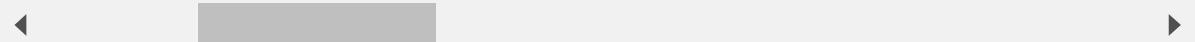
In [134]:

```
1 from sklearn.decomposition import PCA
2 # Extract features with high VIF
3 high_vif_features1 = ['Identification.ID', 'Identification.Model Year']
4
5
6 # Apply PCA
7 pca = PCA(n_components=1) # Reduce to a single component
8 X_pca = pca.fit_transform(df[high_vif_features1])
9
10 # Create a new column in the DataFrame for the PCA component
11 df['PCA_ID_ModelYear'] = X_pca
12
13 # Drop the original correlated features
14 df = df.drop(columns=high_vif_features1)
15
16 print("Transformed DataFrame with PCA feature added:")
17 df.head()
18
```

Transformed DataFrame with PCA feature added:

Out[134]:

Information.Number of Forward Gears	Engine Information.Transmission	Fuel Information.City mpg	Identification.Make	Identification.Ye
0.666667	0.240283	0.476190	0.424175	0
0.666667	0.240283	0.666667	0.424175	0
0.666667	0.489321	0.619048	0.424175	0
0.666667	0.240283	0.619048	0.424175	0
0.666667	0.489321	0.380952	0.424175	0



```
In [135]: 1 #vif
2 from statsmodels.stats.outliers_influence import variance_inflation_factor
3 selected_features = ['PCA_EngineType_highway',
4                      'PCA_ID_ModelYear',
5                      'Fuel Information.Fuel Type_E85']
6
7 X = df[selected_features]
8
9 # Create a DataFrame to store VIF values
10 vif_data = pd.DataFrame()
11 vif_data["Feature"] = X.columns
12 vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(
13 len(X.columns))]
14 print(vif_data)
```

	Feature	VIF
0	PCA_EngineType_highway	23.623078
1	PCA_ID_ModelYear	23.140748
2	Fuel Information.Fuel Type_E85	1.208917

```
In [ ]: 1 #Now again do pca as 1st 2 have high VIF values
```

In [136]:

```

1 from sklearn.decomposition import PCA
2 # Extract features with high VIF
3 high_vif_features2 = ['PCA_EngineType_highway', 'PCA_ID_ModelYear']
4
5
6 # Apply PCA
7 pca = PCA(n_components=1) # Reduce to a single component
8 X_pca = pca.fit_transform(df[high_vif_features2])
9
10 # Create a new column in the DataFrame for the PCA component
11 df['PCA_final'] = X_pca
12
13 # Drop the original correlated features
14 df = df.drop(columns=high_vif_features2)
15
16 print("Transformed DataFrame with PCA feature added:")
17 df.head()
18

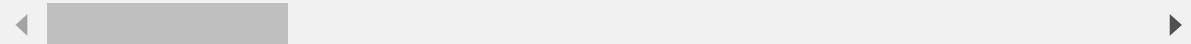
```

Transformed DataFrame with PCA feature added:

Out[136]:

	Dimensions.Height	Dimensions.Length	Dimensions.Width	Engine Information.Number of Forward Gears	Information.Tr...
0	0.545455	0.557312	0.794466	0.666667	
1	0.545455	0.557312	0.794466	0.666667	
2	0.545455	0.557312	0.794466	0.666667	
3	0.545455	0.557312	0.794466	0.666667	
5	0.351779	0.059289	0.241107	0.666667	

5 rows × 21 columns



In [137]:

```

1 #vif
2 from statsmodels.stats.outliers_influence import variance_inflation_factor
3 selected_features = ['Fuel Information.Fuel Type_E85',
4                      'PCA_final']
5
6 X = df[selected_features]
7
8 # Create a DataFrame to store VIF values
9 vif_data = pd.DataFrame()
10 vif_data["Feature"] = X.columns
11 vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(
12 len(X.columns))]
13 print(vif_data)

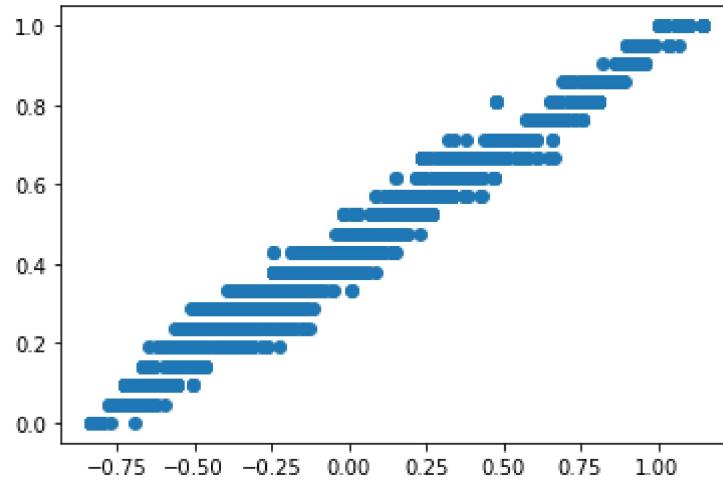
```

	Feature	VIF
0	Fuel Information.Fuel Type_E85	1.192972
1	PCA_final	1.192972

```
In [ ]: 1 #Since VIF is near to 1, it shows Less multicollinearity(independence)
         2 #Now check the Linearity
```

```
In [142]: 1 x=df['PCA_final']
           2 y=df['Fuel Information.City mpg']
           3 plt.scatter(x,y)
```

Out[142]: <matplotlib.collections.PathCollection at 0x2a8e7987b50>



```
In [144]: 1 from sklearn.linear_model import LinearRegression
           2 import matplotlib.pyplot as plt
           3 import seaborn as sns
           4
           5 model=LinearRegression()
           6
           7 X=df[['PCA_final','Fuel Information.Fuel Type_E85']]
           8 Y=df[['Fuel Information.City mpg']]
           9
          10 model.fit(X,Y)
          11 pred=model.predict(X)
```

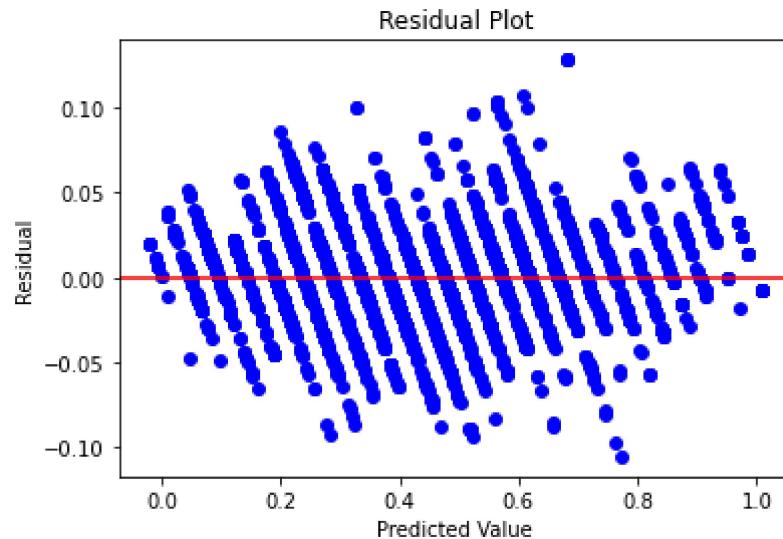
```
In [145]: 1 from sklearn.metrics import r2_score
           2 r2_score_value = r2_score(Y, pred)
           3 print(r2_score_value) #high r^2 means better model
```

0.9814751698519412

```
In [ ]: 1 #checking for homoscedasticity
```

```
In [146]: 1 error=Y-pred
2 plt.scatter(pred,error,color='b')
3 plt.xlabel('Predicted Value')
4 plt.ylabel('Residual')
5 plt.title('Residual Plot')
6 plt.axhline(y=0,color='r')
7
```

```
Out[146]: <matplotlib.lines.Line2D at 0x2a8e78f4640>
```



Conclusion

Hence multiple linear regression is performed with Fuel Information.City mpg as the target variable. Assumptions of Linear Regression model like multicollinearity, linearity and homoscedasticity are verified and a model is trained.

```
In [ ]: 1
```

```
In [ ]: 1
```